

Poisoning the Threshold-Calibration Stage in Federated IoT Anomaly Detection: A Policy-Differentiated Vulnerability Analysis

Salaheddine Nasloubi¹, Mohamed Lachgar², and Hamid Hrimech¹

¹ LAMSAD Laboratory, National School of Applied Sciences (ENSA), Hassan First University, Berrechid, Morocco

² L2IS Laboratory, Higher Normal School (ENS), University Cadi Ayyad, Marrakech, Morocco

Abstract. Federated learning (FL) defenses for IoT anomaly detection focus on the training phase—Byzantine-robust aggregation, gradient filtering, differential privacy. The post-training *threshold-calibration stage*, where each client autonomously collects a benign calibration set and derives its detection threshold, is structurally outside the scope of these defenses. The study characterizes the vulnerability of this stage under a score-level contamination model: a single compromised client replaces a fraction f of its calibration buffer with tail-sampled benign reconstruction scores, shifting the detection threshold without altering model weights, training data, or test labels. Across three threshold-policy families (GLOBAL, LOCAL, CLUSTER) evaluated on the N-BaIoT nine-device federation, the score-level raising attack achieves a 100% pass rate on the pre-specified materiality gate (Gate-1; §4) at all non-zero fractions, degrading victim true-positive rates by 2.4 to 7.9 percentage points across injection fractions (4.9–7.9 pp at $f=0.4$). The score-level lowering attack achieves elevated false-alarm rates for LOCAL and CLUSTER, increasing the cross-client false-positive coefficient of variation $\Delta CV(\text{FPR})$ by 0.05–0.06 (a lower bound under REPLACE-FIXED-BUDGET); the $\Delta \text{TPR} > 0$ reflects threshold displacement, not a security benefit, as it increases the alarm burden that may lead operators to disable the detector. The lowering attack is structurally bounded for GLOBAL due to $1/N$ dilution. Policy choice determines not just per-victim magnitude but threshold-propagation scope: LOCAL confines harm to the compromised client; CLUSTER can transmit small threshold perturbations through co-cluster structure; GLOBAL dilutes per-victim shift but propagates it to all federation members. The reported results are at the score-substitution level: the analysis characterizes a necessary calibration-channel vulnerability mechanism, not a complete traffic-generation attack, and end-to-end traffic-level realization is future work.

Keywords: Federated Learning · Anomaly Detection · IoT Security · Poisoning Attack · Threshold Calibration · Vulnerability Analysis

1 Introduction

Federated learning (FL) is commonly used for IoT anomaly detection: devices contribute to a shared model without exposing raw traffic, and personalized detection thresholds let each node adapt to its own traffic distribution [15,18]. The dominant security concern in federated anomaly detection focuses on the *training phase*—Byzantine clients injecting adversarial gradients, label-flipping poisoning, or model-inversion attacks [5,4,20]. Aggregation-phase defenses such as FLTrust, Krum, and Trimmed-Mean target exactly this surface [8,6,21].

A post-training coverage gap. The *threshold-calibration stage* is the post-training, pre-deployment step in which each client collects a window of benign traffic, computes local reconstruction scores, and derives its detection threshold at a fixed quantile q . Training- and aggregation-phase defenses studied in the FL-IDS literature do not monitor this post-training calibration stage: they complete before calibration begins and cannot observe it. If a single client’s benign calibration buffer can be contaminated with scores derived from benign calibration records that are anomalously high (or low) relative to that client’s own distribution, the resulting threshold shifts without any change to model weights, training data, labels, or gradients. No claim is made that a dedicated calibration-stage defense is impossible—only that the training- and aggregation-phase defenses studied in the FL-IDS literature do not cover this stage.

Study focus. The study identifies threshold calibration as an under-analyzed attack surface and characterizes its vulnerability under a score-level contamination model. Concretely, an adversary with access to a single client’s calibration buffer replaces a fraction f of entries with tail-sampled scores drawn from the same client’s benign distribution (*calibration-stage poisoning*). The evaluation covers three threshold-policy families—GLOBAL (federation-wide mean), LOCAL (per-client), and CLUSTER (cluster-averaged)—across all nine physical-device clients in the N-BaIoT dataset [17] at four injection fractions.

To address the objection—“this is just data poisoning applied to calibration data”—calibration-stage poisoning is structurally distinct along three axes (Table 1): it is *temporally isolated* (strictly after model convergence, bypassing training-phase defenses); it exploits a *phase-specific trust model* (calibration data is collected autonomously with no central validation gate); and its effects are *policy-dependent* (the same injection fraction yields qualitatively different threshold shifts and threshold propagation depending on whether thresholds aggregate globally, locally, or by cluster).

Contributions.

1. The analysis isolates the post-training threshold-calibration stage as an attack surface not covered by training-phase and aggregation-phase defenses, distinct from training-data and aggregation-phase poisoning (§3).
2. The evaluation presents a controlled vulnerability characterization under score-level calibration contamination across three threshold policies, showing that the score-level raising attack reliably degrades victim detection rates by 2.4 to 7.9 percentage points across non-zero injection fractions (4.9–7.9 pp at $f=0.4$) with 95% bootstrap confidence intervals excluding zero (§5).

3. The results document policy-differentiated behavior: LOCAL has the largest per-victim threshold shifts with harm confined to the victim (no cross-client threshold propagation); CLUSTER has potential intra-cluster threshold perturbation; downstream non-victim harm is not separately measured; GLOBAL has magnitude-diluted fleet-wide propagation; the lowering attack succeeds for LOCAL and CLUSTER but is conditionally bounded for GLOBAL (§5).

Scope. All claims are bounded to the nine-device N-BaIoT setting with autoencoder detectors and FedAvg aggregation, and the poisoning mechanism operates at the score level; end-to-end traffic-generation realizability is future work (§6). No claim is made about deployability, evasion guarantees, or generalization beyond this setting. Defense design against calibration-stage poisoning is out of scope, though §6 discusses properties a defense would need to satisfy.

2 Background and Related Work

Federated anomaly detection for IoT. Network-level anomaly detection using federated learning allows devices to collaborate on a shared model while keeping raw traffic local [18]. Autoencoder-based reconstruction-error detectors are a standard approach: each client trains a local encoder-decoder on benign traffic; anomalies are flagged when reconstruction error exceeds a per-client detection threshold [22]. The N-BaIoT benchmark [17] provides labeled benign and Mirai/BASHLITE traffic for nine physical devices spanning five device types, establishing a widely used evaluation platform for FL-IDS research.

Defenses. Existing defenses target the *training phase*. Byzantine-robust rules, including Krum [6], Trimmed-Mean, and Median [21], filter or downweight outlier gradients before aggregation. FLTrust [8] anchors aggregation to a trusted server dataset. Differential privacy approaches add noise to local updates to bound update privacy leakage [11]; DP is a privacy mechanism for the training/update channel, not a poisoning defense, and this study does not treat it as one. These defenses operate on gradients or model updates; they complete before the threshold-calibration stage begins.

Threshold calibration and poisoning. By construction, a detection threshold at fixed quantile q is sensitive to the scores admitted to the calibration distribution: tail-high contamination raises the boundary (missed detections); tail-low contamination lowers it (false alarms). Poisoning attacks on anomaly detectors have been studied in centralized settings [13,9,3], e.g. Kloft & Laskov’s [13] online poisoning via a continuous training stream; these assume a trusted aggregator or persistent model-update access, both absent from FL’s post-training calibration stage, where calibration data never leaves the client and no server-side sanitization is possible. Prior work on FL model poisoning [2] and FL inference attacks [10] focuses on the gradient/weight surface. The literature reviewed for this study does not identify prior work that isolates post-training threshold calibration as the attacked channel in federated threshold-based IoT anomaly detection under otherwise-unchanged training, model, aggregation, and

Table 1. Positioning relative to adjacent poisoning literature. DATP-CP differs by acting strictly after training, on the autonomously collected benign calibration buffer, leaving model weights, gradients, training data, and test data unchanged.

Prior work type	Attack surface	What changes	Distinction from DATP-CP
Training-data poisoning [5,20]	Training set / labels	Model parameters	DATP-CP leaves training data and model unchanged
FL model poisoning [4,2]	Client gradients	Model weights	Needs no gradient access; acts after convergence
Byzantine aggregation [6,21]	Aggregation inputs	Server aggregate	Does not touch aggregation; harm is post-aggregation
Centralized AD poisoning [13,9]	Central train stream	Decision boundary	No trusted sanitizer exists in FL calibration
Threshold personalization [12,19]	Threshold policy	Per-client/cluster τ	Studies utility, not adversarial contamination
DATP-CP	Calibration scores	Threshold τ only	Score-level, post-training, single-client

test artifacts (Table 1). No claim is made to a new class of poisoning in the abstract sense; the contribution is isolating the post-training calibration channel in federated threshold-based anomaly detection and showing policy-dependent victim effects and threshold propagation under a controlled score-level intervention. The closest neighbors treat the calibration/thresholding pipeline as trusted: Laridi et al. [14] propose summary-statistics-based federated thresholding without adversarial contamination, and Sáez-de-Cámara et al. [7] propose a clustered federated IoT anomaly-detection architecture assuming a benign calibration pipeline; DATP-CP instead treats client-local benign calibration as the attacked channel.

Threshold policies. Federated threshold personalization spans GLOBAL, LOCAL, and CLUSTER designs [12,19] (defined precisely in §4); their security implications have not been studied under calibration-stage adversaries.

3 Threat Model

Attack surface. The threshold-calibration stage occurs after global model convergence and before threshold-dependent detection begins. Each client i collects a benign calibration set $\mathcal{C}_i = \{s_1, \dots, s_{n_i}\}$ of reconstruction scores from its own traffic, then sets its local threshold $\tau_i = \text{quantile}(\mathcal{C}_i, q)$ at a fixed quantile q (here $q=0.95$). Depending on the federation’s threshold policy, τ_i may be used directly (LOCAL), averaged across clients (GLOBAL), or averaged within a calibration-score cluster (CLUSTER).

No existing FL training-phase or aggregation-phase defense monitors or validates the calibration buffer: it is collected autonomously by the client after the aggregation round completes.

Adversary model. The adversary controls a single compromised client v , called the *poisoned client*. Under LOCAL policy, only client v ’s own detection degrades; under GLOBAL policy, all federation members receive the contaminated threshold. The term “victim” refers to client v itself when measuring per-client detection harm, and “non-victim clients” refers to uninvolved federation members.

The adversary has *gray-box score-level access*: since the adversary controls the client, it can read and modify only its own local calibration buffer before threshold computation. It cannot modify training data, training labels, test scores, test labels, model weights, gradients, the aggregation procedure, other clients’ data, client eligibility, or the threshold-policy definition (GLOBAL/LOCAL/CLUSTER are fixed by the federation). The reservoir is drawn exclusively from the same client’s own benign calibration split (§4): attack-labeled samples, test scores, and training scores never enter it, by construction. This is a strictly weaker capability than model-poisoning adversaries, which require gradient injection. The adversary operates at the score level; generating traffic that achieves a target reconstruction score is out of scope (end-to-end realizability is discussed in §6). The REPLACE-FIXED-BUDGET mechanism samples with replacement from the tail reservoir, which at high fractions produces duplicate scores detectable by a uniqueness check (disclosed in §6).

Attack mechanics. The evaluation considers two attack objectives:

Raise. THRESHOLD-RAISE (primary): $f=0$ is the unpoisoned clean baseline ($m=0$, no replacement); for $f>0$, replace $m = \max(1, \text{round}(fn_v))$ positions in $\mathcal{C}_v$ with scores drawn with replacement from the top-10% tail of $\mathcal{C}_v$ (HIGH-SCORE-BENIGN source). All replacement scores are benign-origin reconstruction scores drawn from the victim’s own clean calibration distribution; AUROC is invariant by construction (test scores and labels are unchanged under poisoning). A raised threshold τ_v increases the detection threshold, suppressing alarms at the victim site.

Lower. THRESHOLD-LOWER (secondary, conditional): replace m positions with scores from the bottom-10% tail (LOW-SCORE-BENIGN). A lowered threshold increases the false-positive rate (alarm burden) and FPR disparity; the resulting $\Delta\text{TPR}>0$ reflects threshold displacement and is a secondary indicator of operating-point shift. This attack is structurally conditional: random position replacement displaces upper-tail support with probability $\approx f$.

Negative control. RANDOM-BENIGN: replacement scores are sampled uniformly from the full calibration distribution rather than a tail. Any effect attributable to score contamination rather than position perturbation should vanish in this condition.

Scope boundary. The analysis characterizes the vulnerability at the score level. A complete end-to-end attack would additionally require generating network traffic whose autoencoder reconstruction score lands in the target tail—this




Fig. 1. Attack surface: the threshold-calibration stage occurs after FL training completes. Green (solid-border) boxes mark the training/aggregation phases covered by existing FL-IDS defenses; the dashed red region marks the calibration stage, which those defenses do not observe (no defense is implemented here). An adversary with score-level access to one client’s calibration buffer can shift τ without modifying model weights, gradients, or labels.

is future work (§6). The score-level characterization is a controlled experiment, not a deployed attack: it isolates a necessary calibration-channel vulnerability mechanism, not a complete traffic-generation attack. Claims are bounded to the single-client, single-compromise scenario; multi-client collusion is not evaluated.

4 Experimental Protocol

4.1 Dataset and Federation

The evaluation uses N-BaIoT [17], a labeled IoT network traffic dataset covering nine physical devices spanning five device types (doorbell, baby monitor, security camera, thermostat, webcam). Each device constitutes one federation client; all nine are eligible for calibration. Calibration set sizes range from 2,622 (Ecobee Thermostat) to 35,048 (Philips Baby Monitor) benign samples per device (all ≥ 100 , satisfying the eligibility threshold). Each device’s benign stream is partitioned chronologically into train/calibration/test segments in a 60/1/20/1/18 ratio (the two 1% gaps separate segments to avoid temporal leakage); training and calibration use benign traffic only, and test data, test labels, and training labels are never touched by the attack. Attack-labeled records are loaded from the N-BaIoT attack folders, scaled with the device scaler fitted on training data, scored as the `test_attack` split, and used only as the positive class for TPR evaluation; they never enter FedAvg training, benign calibration arrays, poisoning reservoirs, threshold recomputation, or injected buffers. The federation trains a per-client autoencoder on benign traffic using FedAvg [16] for feature extraction, then enters the threshold-calibration stage at a fixed quantile $q=0.95$. Table 2 lists the autoencoder architecture, training hyperparameters, and the calibration

Table 2. Experimental configuration for reproducibility.

Component	Configuration
<i>Autoencoder architecture</i>	
Input features	115 (N-BaIoT statistical traffic features)
Encoder layers	115 $\rightarrow$ 80 $\rightarrow$ 40 $\rightarrow$ 20; decoder mirrors
Activation	ReLU (no batch normalization)
Loss	Mean squared reconstruction error
Optimizer	Adam, lr = 0.001, weight_decay = 0
<i>Federated training (FedAvg)</i>	
Local epochs	1
Aggregation	Sample-count-weighted mean
Rounds	40 initial, 150 max (conv.: $\Delta\text{loss} \leq 0.005$ over 10)
Batch size	256
<i>Calibration (per-device sizes in released split manifest)</i>	
Benign samples / device	2,622 (Ecobee) – 35,048 (Philips)
Threshold quantile q	0.95

size range; injection fractions, seeds, and the reservoir are detailed under Attack Parameterization below.

4.2 Threshold Policies

Three threshold-policy families are evaluated, covering the main design points in the personalization literature:

Global.

$$\tau_i^{\text{local}} = \text{quantile}(\mathcal{C}_i, q), \quad \tau^{\text{global}} = \frac{1}{N} \sum_{i=1}^N \tau_i^{\text{local}}.$$

Global assigns the same τ^{global} to every client; a single poisoned local threshold contributes only $1/N$ of the shared threshold shift.

Local. Each client retains $\tau_i = \tau_i^{\text{local}}$. Shifts are confined to the compromised client; no cross-client threshold propagation.

Cluster. Clients are grouped on a standard-scaled 4-feature fingerprint (mean, std, skew, p_{95} of $\mathcal{C}_i$) via k -means++ [1] ($K=3$, `random_state=42`; clustering recomputed after each attack draw so a poisoned fingerprint can change the compromised client’s cluster assignment). Within each cluster, the shared threshold is the cluster mean; the higher seed-to-seed variance in `CLUSTER CV(FPR)` reflects this cluster-assignment sensitivity. A compromised client may perturb the threshold for co-cluster members (*intra-cluster threshold perturbation*); clients in other clusters are unaffected.

4.3 Attack Parameterization

For each (policy, attack objective, injection fraction $f \in \{0.0, 0.1, 0.2, 0.4\}$) combination, the protocol runs 10 paired seeds: each seed selects a training configuration and an independent poisoning draw. A clean copy of the calibration array

is preserved; only the poisoned copy is used for threshold computation. The four source-objective pairs are: (1) HIGH-SCORE-BENIGN + THRESHOLD-RAISE (primary raising attack), (2) LOW-SCORE-BENIGN + THRESHOLD-LOWER (primary lowering attack), (3) RANDOM-BENIGN + THRESHOLD-RAISE (objective-matched negative control for the raising attack), and (4) RANDOM-BENIGN paired with THRESHOLD-LOWER (negative control for the lowering attack). The attack is applied to each of the nine clients as victim in turn (single-client scope), yielding $4 \text{ source-objective pairs} \times 3 \text{ threshold policies} \times 4 \text{ injection fractions} \times 9 \text{ victim clients} \times 10 \text{ seeds} = 4,320 \text{ experimental cells total}$. Training seed $i \in \{0, \dots, 9\}$ is paired with poisoning seed $i+100$ and analysis seed $i+300$, giving statistically independent stochasticity across training, calibration sampling, and bootstrap resampling.

4.4 Statistical Analysis

Gate 1 (material shift, primary gate). Threshold shift $\Delta\tau_v = \tau_v^{\text{poisoned}} - \tau_v^{\text{clean}}$ is computed per seed. A seed is *Gate-1 significant* when $|\Delta\tau_v| > \delta\tau_v$ and the sign is consistent with the attack objective, where the materiality threshold is $\delta\tau_v = \max(0.1 \times \text{IQR}(\mathcal{C}_v^{\text{clean}}), 0.01 \times \text{median}_j(\text{IQR}(\mathcal{C}_j^{\text{clean}})))$. Gate-1 passes when $\geq 8/10$ seeds are significant and the victim-majority condition ($\geq 5/9$ victims significant, i.e. a simple majority of the nine federation clients) holds.

Gate 2 (directional excess over control). Bootstrap 95% CI on mean $\Delta\tau$ excludes zero for the primary strategy but spans zero for the negative control (RANDOM-BENIGN).

Gate 3 (downstream harm). Mean victim TPR change, ΔTPR , is directionally consistent (negative for raising, positive for lowering) across $\geq 8/10$ seeds.

Confidence intervals. All CIs are 95% percentile bootstrap over 10 seed-level aggregates, with $B=10,000$ resamples. Primary claim gates are evaluated over these paired seed-level aggregates; the per-victim-seed cells are descriptive support, not independent units. With ten paired seeds, percentile-bootstrap intervals are approximate uncertainty estimates rather than exact coverage guarantees, and may deviate from nominal at the tails.

Gate constant justification. The 10% IQR coefficient for Gate-1 materiality ($0.1 \times \text{IQR}$) exceeds the calibration instability observed in the clean baseline while remaining below the expected tail-injection shift at $f=0.1$; the 1% cross-client floor ($0.01 \times \text{median}_j(\text{IQR})$) prevents degenerate clients with near-zero score spread from counting as material on arbitrarily small shifts. Both constants were locked from clean-artifact statistics before any poisoned run was executed.

4.5 Isolation Sanity Check

To confirm that the attack is confined to the calibration stage, the protocol verifies: (i) AUROC is invariant under poisoning for all cells, (ii) per-client rank

ordering of clean anomaly scores is unchanged, and (iii) model weights before and after calibration are bitwise identical. All three hold in all 4320 cells.

5 Results

5.1 Clean-Baseline Policy Profile

Table 3 reports clean-baseline dispersion statistics across all ten seeds. $CV(FPR)$, the coefficient of variation of per-client false-positive rates ($ddof=0$), is lower when FPR treatment is more uniform across the federation. LOCAL achieves the lowest $CV(FPR)$ (0.333 ± 0.034). GLOBAL exhibits the highest $CV(FPR)$ (0.905 ± 0.039) because a single federation-wide threshold cannot simultaneously fit nine devices with different traffic distributions. CLUSTER lies in between (0.615 ± 0.143) with higher variance across seeds, reflecting sensitivity to cluster assignment. Bootstrap 95% CIs on the pairwise $CV(FPR)$ difference exclude zero: Global–Local = 0.572 [0.537, 0.607]; Global–Cluster = 0.291 [0.214, 0.358]; Cluster–Local = 0.281 [0.201, 0.374], confirming that all three policies are statistically distinguishable on baseline fairness. Figure 2 shows the per-client FPR distributions for GLOBAL and LOCAL under the clean baseline. This fairness gain is not a uniform improvement: LOCAL improves cross-client FPR equity but does not dominate all accuracy-oriented metrics. In particular, GLOBAL retains the best P10 Macro-F1 (0.398 vs. LOCAL 0.335, Table 3), so the worst-tail client is detected better under the federation-mean threshold. The evaluation therefore reports both fairness-style dispersion ($CV(FPR)$) and detection-quality (P10 Macro-F1) metrics, plus the underlying absolute mean FPR/TPR, rather than a single summary.

Table 3. Clean-baseline dispersion and detection statistics per threshold policy (N-BaIoT, $n=10$ seeds, 9 clients). CV: coefficient of variation ($ddof=0$; lower is more uniform); BA: balanced accuracy; FPR/TPR are client/seed means. Bold = best per column. * Actual std ≈ 0.0002 ; rounds to zero.

Policy	CV(FPR)	CV(TPR)	Worst BA	P10 Macro-F1	Mean FPR	Mean TPR
GLOBAL	0.905 ± 0.039	0.255 ± 0.030	0.657 ± 0.008	0.398 ± 0.053	0.041	0.793
LOCAL	0.333 ± 0.034	0.296 ± 0.036	0.655 ± 0.005	0.335 ± 0.058	0.043	0.777
CLUSTER	0.615 ± 0.143	0.313 ± 0.009	0.656 ± 0.004	$0.299 \pm 0.000^*$	0.043	0.742

5.2 Raising Attack: Primary Results

Table 4 and Figure 3(a) report threshold-shift and downstream detection harm for the HIGH-SCORE-BENIGN + THRESHOLD-RAISE strategy.

All three policies pass all three gates. The seed-level Gate-1 pass rate is 100% (10/10 seeds) at every non-zero fraction for all three policies. For LOCAL and CLUSTER, all nine clients are individually significant in every seed. For GLOBAL at $f=0.1$, two devices (Provision PT-737E and PT-838) are intermittently non-significant (2/10 and 4/10 seeds respectively), but the victim-majority condition ($\geq 5/9$ per seed) still holds because the remaining seven clients




Fig. 2. Per-client FPR under clean baseline, GLOBAL vs. LOCAL threshold (seed 0; CV(FPR) values across all 10 seeds reported in Table 3). GLOBAL forces a single threshold on nine devices with heterogeneous score distributions, producing CV(FPR) = 0.87 for seed 0 (mean = 0.905 ± 0.039 across seeds). LOCAL fits each client’s own calibration distribution, achieving CV(FPR) = 0.33 on average. CLUSTER is omitted here; its baseline is in Table 3 (CV(FPR) = 0.615 ± 0.143).

Table 4. Threshold shift and detection harm under calibration-stage poisoning (N-BaIoT, $n=10$ seeds). $\Delta\tau$: mean threshold shift [95 % CI]. G1: fraction of 10 seeds passing Gate-1. ΔTPR (pp): victim detection-rate change; negative for Raise (harm), positive for Lower (alarm-burden increase, not a security benefit). Lowering rows ($\dagger$): $\Delta\text{CV}(\text{FPR})$ is the primary harm metric (“—” = not the primary axis for Raise/Neg. Ctrl). NEG. CONTROL: RANDOM-BENIGN at $f=0.4$, similarly near-null at all fractions; bold=zero G1. Absolute FPR/TPR at $f=0.4$ are given in §5.

Policy	Attack	f	$\Delta\tau$ [95 % CI]	G1	ΔTPR (pp)	$\Delta\text{CV}(\text{FPR})$
GLOBAL	Raise	0.1	+0.033 [+0.031, +0.035]	100 %	−2.4	—
		0.2	+0.062 [+0.059, +0.065]	100 %	−4.6	—
		0.4	+0.100 [+0.096, +0.105]	100 %	−6.1	—
LOCAL	Raise	0.1	+0.298 [+0.281, +0.314]	100 %	−6.0	—
		0.2	+0.558 [+0.535, +0.581]	100 %	−7.3	—
		0.4	+0.902 [+0.864, +0.944]	100 %	−7.9	—
CLUSTER	Raise	0.1	+0.253 [+0.231, +0.277]	100 %	−3.9	—
		0.2	+0.504 [+0.475, +0.532]	100 %	−4.9	—
		0.4	+0.899 [+0.859, +0.941]	100 %	−4.9	—
GLOBAL	Lower $\dagger$	0.1	−0.006 [−0.006, −0.006]	60 %	+0.8	−0.007
		0.2	−0.010 [−0.011, −0.010]	100 %	+1.1	−0.011
		0.4	−0.018 [−0.020, −0.017]	100 %	+2.1	−0.019
LOCAL	Lower $\dagger$	0.1	−0.053 [−0.057, −0.049]	100 %	+2.8	−0.002
		0.2	−0.094 [−0.100, −0.088]	100 %	+5.9	+0.005
		0.4	−0.164 [−0.176, −0.154]	100 %	+11.8	+0.053
CLUSTER	Lower $\dagger$	0.1	−0.047 [−0.065, −0.031]	100 %	+2.2	+0.007
		0.2	−0.100 [−0.115, −0.084]	100 %	+5.4	+0.011
		0.4	−0.154 [−0.167, −0.142]	100 %	+10.3	+0.060
GLOBAL	Neg. Ctrl	0.4	−0.000 [−0.001, +0.001]	0 %	+0.3	—
LOCAL	Neg. Ctrl	0.4	−0.002 [−0.009, +0.006]	0 %	+0.9	—
CLUSTER	Neg. Ctrl	0.4	−0.009 [−0.022, +0.003]	0 %	+0.4	—

are consistently significant. Bootstrap CIs on $\Delta\tau$ exclude zero and remain non-overlapping with the RANDOM-BENIGN control, which produces CIs spanning zero, a 0% Gate-1 pass rate, and $|\Delta\text{TPR}| \leq 1$ pp at every fraction and policy




Fig. 3. Mean threshold shift $\Delta\tau$ vs. injection fraction f for the raising (a) and lowering (b) attacks across three policies. Error bars: 95% bootstrap CI over 10 seeds. GLOBAL exhibits the smallest magnitude in both directions due to $1/N$ federation averaging. RANDOM-BENIGN negative control (not shown) produces CIs spanning zero and $G1=0\%$ at all fractions and policies; representative values at $f=0.4$ (Neg. Ctrl rows).

(Table 4, Neg. Ctrl rows, reports representative $f=0.4$ values; all fractions are similarly near-null and the complete matrix is archived in the artifact bundle). The targeted effects are therefore attributable to tail contamination, not to position perturbation alone.

Policy magnitude. Raising-attack magnitude differs by policy. LOCAL and CLUSTER exhibit the largest threshold shifts: at $f=0.4$, $\Delta\tau = +0.902$ $[+0.864, +0.944]$ and $+0.899$ $[+0.859, +0.941]$ respectively. GLOBAL shows a structurally diluted but still significant shift: $+0.100$ $[+0.096, +0.105]$ at $f=0.4$. The dilution is expected by construction—GLOBAL’s threshold is the federation mean, so a single-client shift is divided by $N=9$ —but the empirical question is whether the residual magnitude ($\approx 10\%$ of the local shift) suffices to produce measurable downstream harm. It does.

Downstream harm (Gate 3). Mean victim ΔTPR is negative and directionally consistent across all seeds for all three policies. At $f=0.4$: LOCAL -7.9 pp, CLUSTER -4.9 pp (plateauing from $f=0.2$), GLOBAL -6.1 pp. Absolute clean→poisoned victim TPR at $f=0.4$ (Table 3 clean column): GLOBAL $0.793 \rightarrow 0.732$, LOCAL $0.777 \rightarrow 0.698$, CLUSTER $0.742 \rightarrow 0.693$. The GLOBAL threshold propagation is notable: all eight non-victims receive the identical $\Delta\tau$ (Table 5); their TPR was not separately recomputed, so downstream non-victim TPR impact is inferred, not measured, and heterogeneity ($\text{CV}(\text{FPR}) = 0.905$) means it could differ from the victim’s. LOCAL/CLUSTER non-victim threshold propagation is bounded by construction.

5.3 Lowering Attack: Secondary Results (Conditional)

Figure 3(b) shows threshold shift for the LOW-SCORE-BENIGN + THRESHOLD-LOWER strategy. The adversary’s goal is to increase the false alarm rate; the $\Delta\text{TPR} > 0$ below is a threshold-displacement indicator, not a security benefit. Elevated alarm burden may lead operators to disable the anomaly detector.

Local and Cluster succeed; Global is marginal. LOCAL achieves $\Delta\tau = -0.164$ $[-0.176, -0.154]$ at $f=0.4$ with elevated FPR disparity ($\Delta\text{CV}(\text{FPR}) = +0.053$); Gate-3 $\Delta\text{TPR} = +11.8$ pp confirms threshold displacement. CLUSTER similarly achieves -0.154 $[-0.167, -0.142]$ at $f=0.4$ with $\Delta\text{CV}(\text{FPR}) = +0.060$ and $\Delta\text{TPR} = +10.3$ pp. GLOBAL is structurally bounded: at $f=0.4$, $\Delta\tau = -0.018$ $[-0.020, -0.017]$, Gate-1 passes only at $f \geq 0.2$, and $\Delta\text{TPR} = +2.1$ pp.

FPR disparity (operational harm). The operational harm metric is $\Delta\text{CV}(\text{FPR})$ (change in coefficient of variation of per-client FPR, dimensionless): $+0.053$ (LOCAL, 10/10 seeds) and $+0.060$ (CLUSTER, 8/10 seeds) at $f=0.4$, while GLOBAL shows -0.019 (shared-threshold proportional FPR rise reduces relative disparity). Absolute clean $\rightarrow$ poisoned FPR at $f=0.4$ (Table 3 clean column): GLOBAL $0.041 \rightarrow 0.043$, LOCAL $0.043 \rightarrow 0.046$, CLUSTER $0.043 \rightarrow 0.049$.

Mechanism and conditionality. Lowering a $q=0.95$ quantile requires displacing scores in the upper-tail support; random position replacement achieves this with probability $\approx f$ (40 % at $f=0.4$), leaving 60 % of upper-tail support intact. The GLOBAL policy compounds this via $1/N$ dilution. These results are presented as a lower bound under REPLACE-FIXED-BUDGET; targeted upper-tail removal would achieve larger shifts.

5.4 Policy-Differentiation Summary

Table 5. Threshold-propagation profile under the raising attack ($f=0.4$, $n=10$ seeds). Blast: mean fraction of non-victims with a statistically significant threshold shift (significance only, not measured downstream harm; non-victim ΔTPR was not separately recomputed for any policy—see §5). [†] Non-victim threshold equals victim’s by construction (shared threshold). [‡] K-means reassignment (churn); mean non-victim $\Delta\tau=0.0003$, std = 0.017.

Policy	Victim $\Delta\tau$	Victim ΔTPR (pp)	Non-vic. $\Delta\tau$	Blast
GLOBAL	+0.100	−6.1	+0.100 [†]	$\approx 8/8$
CLUSTER	+0.899	−4.9	$\approx 0.000^{\ddagger}$	$\approx 6/8$
LOCAL	+0.902	−7.9	0.000	0/8

Table 5 quantifies the threshold-propagation contrast under the raising attack at $f=0.4$. GLOBAL propagates the threshold shift to all nine clients: the non-victim $\Delta\tau$ equals the victim $\Delta\tau$ by construction (shared threshold), and blast encompasses $\approx 8/8$ non-victims in nearly every seed. CLUSTER achieves near-LOCAL victim-level harm; non-victim threshold changes are small in magnitude (mean $\Delta\tau \approx 0.0003$, std = 0.017, driven by K-means cluster reassignment rather than aggregation), with 6/8 non-victims showing statistically significant but small-magnitude threshold perturbations. The blast count measures threshold significance, not downstream harm: non-victim rates were not recomputed for CLUSTER; since this $\Delta\tau$ is roughly two orders smaller than the $f=0.1$ raising shift (−2.4 to −6.0 pp victim harm, Table 4), large non-victim ΔTPR changes are implausible but unverified—unlike GLOBAL, which transmits $\Delta\tau$ to every non-victim by construction (also inferred, not measured). LOCAL yields the highest per-victim shift with zero non-victim blast by policy design.

Policy design trades isolated per-victim exposure (LOCAL) against diluted fleet-wide propagation (GLOBAL); CLUSTER carries churn-level threshold perturbations without direct aggregation propagation.

6 Discussion and Limitations

Score-level proxy and detectability. The experiment operates at the score level: reconstruction scores in the calibration buffer are replaced programmatically rather than through generated traffic. A complete end-to-end attack would additionally require generating network packets whose autoencoder reconstruction error lands in the target tail of the victim’s score distribution—a non-trivial requirement spanning feature engineering, the encoder-decoder nonlinearity, and physical network behavior. This work is presented as a *vulnerability-surface characterization*, not a deployed attack; traffic-level realizability is future work. A related implementation limitation: with-replacement sampling from the tail reservoir at $f=0.4$ produces duplicate scores in the calibration buffer—for the smallest per-client calibration set ($n_v \approx 2,622$), on the order of 10^3 expected collision pairs, easily flagged by a uniqueness check; evading this requires score perturbation at the cost of leaving the exact tail boundary. Detectability remains an open problem.

GLOBAL threshold propagation, lowering bounds, and policy implications. The $1/N$ -diluted GLOBAL shift is definitional; what is empirical is that the residual magnitude ($\approx 10\%$ of the local shift, $\Delta\text{TPR} = -6.1$ pp at $f=0.4$) remains measurable above the seed-to-seed variance. The Local/Cluster lowering results are a *lower bound* under REPLACE-FIXED-BUDGET (§5): adversaries with access to the buffer’s rank structure could achieve larger threshold displacements and alarm-burden effects. These results reveal a structural trade-off in policy design: LOCAL maximizes per-client isolation but maximizes per-victim exposure, whereas GLOBAL limits per-victim shift via averaging but propagates residual threshold shifts federation-wide; defense designers must choose which risk to bound.

Scope limitations. Claims are bounded to N-BaIoT (9 clients; generalizability to other federated IoT datasets and larger federation sizes is open), autoencoder architecture, FedAvg, and $q=0.95$ (quantile sensitivity is future work). The experiment assumes a single compromised client; multi-client collusion is a direct extension. Generalizability to other anomaly detectors, FL algorithms, and heterogeneous federation profiles is not established. The $n=10$ seed bootstrap CI coverage may differ from nominal at the tails. Defense design for the calibration channel (threshold validation, calibration-set integrity constraints) is out of scope and merits independent study.

7 Conclusion

The study characterized the post-training threshold-calibration stage in federated IoT anomaly detection as an attack surface that training-phase and

aggregation-phase defenses do not address. Under a score-level contamination model on the N-BaIoT nine-device federation (FedAvg autoencoder, $q=0.95$, single compromised client), the raising attack degrades victim detection rates by 2.4–7.9 pp across non-zero injection fractions (4.9–7.9 pp at $f=0.4$) with 100 % Gate-1 pass rates and 95 % bootstrap CIs excluding zero; these intervals are descriptive uncertainty estimates over ten paired seeds. The lowering attack shifts thresholds downward, raising the alarm-burden indicator ΔTPR by 10–12 pp—an operating-point displacement, not a security benefit—and $\Delta\text{CV}(\text{FPR})$, the primary lowering-harm metric, by 0.05–0.06 (lower bound under REPLACE-FIXED-BUDGET) for LOCAL and CLUSTER. The policy ordering—per-victim exposure maximized under LOCAL, diluted but federation-wide threshold propagation under GLOBAL, small intra-cluster threshold perturbations under CLUSTER—implies calibration-stage hardening should be matched to the aggregation design, not only the training-phase defense. These findings isolate a necessary calibration-channel vulnerability mechanism, not a complete traffic-generation attack. Future work should address traffic-level realizability, calibration-stage integrity defenses, and generalization across detector architectures and datasets.

Reproducibility. All experiments are pipeline-generated and locally archived (code commit 396ea81). The artifact bundle comprises the experiment configuration; seed scheme (training/poisoning/analysis seeds paired per run); split manifest with per-device calibration sizes; per-cell results; negative-control records; isolation sanity checks; and figure/table scripts. No poisoned metric was used for model checkpoint or hyperparameter selection. Code and artifacts will be released publicly on acceptance and are available to reviewers on request.

Disclosure of Interests. The authors have no competing interests to declare that are relevant to the content of this article.

References

1. Arthur, D., Vassilvitskii, S.: k-means++: The advantages of careful seeding. In: Proceedings of the 18th Annual ACM-SIAM Symposium on Discrete Algorithms (SODA). pp. 1027–1035 (2007)
2. Bagdasaryan, E., Veit, A., Hua, Y., Estrin, D., Shmatikov, V.: How to backdoor federated learning. In: Proceedings of the 23rd International Conference on Artificial Intelligence and Statistics (AISTATS). pp. 2938–2948. PMLR (2020)
3. Barreno, M., Nelson, B., Joseph, A.D., Tygar, J.D.: The security of machine learning. *Machine Learning* **81**(2), 121–148 (2010)
4. Bhagoji, A.N., Chakraborty, S., Mittal, P., Calo, S.: Analyzing federated learning through an adversarial lens. In: Proceedings of the 36th International Conference on Machine Learning (ICML). pp. 634–643. PMLR (2019)
5. Biggio, B., Nelson, B., Laskov, P.: Poisoning attacks against support vector machines. In: Proceedings of the 29th International Conference on Machine Learning (ICML). pp. 1467–1474 (2012)
6. Blanchard, P., El Mhamdi, E.M., Guerraoui, R., Stainer, J.: Machine learning with adversaries: Byzantine tolerant gradient descent. In: Advances in Neural Information Processing Systems (NeurIPS). vol. 30, pp. 119–129 (2017)

7. Sáez-de Cámara, X., Flores, J.L., Arellano, C., Urbieto, A., Zurutuza, U.: Clustered federated learning architecture for network anomaly detection in large scale heterogeneous IoT networks. *Computers & Security* **131**, 103299 (2023)
8. Cao, X., Fang, M., Liu, J., Gong, N.Z.: FLTrust: Byzantine-robust federated learning via trust bootstrapping. In: *Network and Distributed System Security Symposium (NDSS)* (2021)
9. Cretu, G.F., Stavrou, A., Locasto, M.E., Stolfo, S.J., Keromytis, A.D.: Casting out demons: Sanitizing training data for anomaly sensors. In: *2008 IEEE Symposium on Security and Privacy*. pp. 81–95. IEEE (2008)
10. Geiping, J., Bauermeister, H., Dröge, H., Moeller, M.: Inverting gradients—how easy is it to break privacy in federated learning? In: *Advances in Neural Information Processing Systems (NeurIPS)*. vol. 33, pp. 16937–16947 (2020)
11. Geyer, R.C., Klein, T., Nabi, M.: Differentially private federated learning: A client level perspective. In: *NeurIPS Workshop on Private Multi-Party Machine Learning* (2017)
12. Ghosh, A., Chung, J., Yin, D., Ramchandran, K.: An efficient framework for clustered federated learning. In: *Advances in Neural Information Processing Systems (NeurIPS)*. vol. 33, pp. 19586–19597 (2020)
13. Kloft, M., Laskov, P.: Online anomaly detection under adversarial impact. In: *Proceedings of the 13th International Conference on Artificial Intelligence and Statistics (AISTATS)*. pp. 405–412. PMLR (2010)
14. Laridi, S., Palmer, G., Tam, K.M.M.: Enhanced federated anomaly detection through autoencoders using summary statistics-based thresholding. *Scientific Reports* **14**(1), 26704 (2024). <https://doi.org/10.1038/s41598-024-76961-2>
15. Li, T., Sahu, A.K., Talwalkar, A., Smith, V.: Federated learning: Challenges, methods, and future directions. *IEEE Signal Processing Magazine* **37**(3), 50–60 (2020)
16. McMahan, B., Moore, E., Ramage, D., Hampson, S., y Arcas, B.A.: Communication-efficient learning of deep networks from decentralized data. In: *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*. pp. 1273–1282. PMLR (2017)
17. Meidan, Y., Bohadana, M., Mathov, Y., Mirsky, Y., Shabtai, A., Breitenbacher, D., Elovici, Y.: N-BaIoT: Network-based detection of IoT botnet attacks using deep autoencoders. *IEEE Pervasive Computing* **17**(3), 12–22 (2018)
18. Rey, V., Sanchez Sanchez, P.M., Huertas Celdrán, A., Bovet, G.: Federated learning for malware detection in IoT devices. *Computer Networks* **204**, 108693 (2022)
19. Sattler, F., Müller, K.R., Samek, W.: Clustered federated learning: Model-agnostic distributed multitask optimization under privacy constraints. *IEEE Transactions on Neural Networks and Learning Systems* **32**(8), 3710–3722 (2021)
20. Tolpegin, V., Truex, S., Gursoy, M.E., Liu, L.: Data poisoning attacks against federated learning systems. In: *25th European Symposium on Research in Computer Security (ESORICS)*. pp. 480–501. Springer (2020)
21. Yin, D., Chen, Y., Kannan, R., Bartlett, P.: Byzantine-robust distributed learning: Towards optimal statistical rates. In: *Proceedings of the 35th International Conference on Machine Learning (ICML)*. pp. 5650–5659. PMLR (2018)
22. Zhou, C., Paffenroth, R.C.: Anomaly detection with robust deep autoencoders. In: *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. pp. 665–674 (2017)